

17-214/514: Agentic Software Development

Agentic Development and Design for Verification

Hammad Ahmad and Claire Le Goues • Wednesday, August 26, 2026 • **Lecture**

Administrivia

- If you have not yet done so, please fill out the intake form by **today at 5pm**.
 - If you filled out the form, you should have received a repo invite for Assignment 1.
- If you have a \$20/month subscription (Claude Pro, etc.), you don't need to preemptively upgrade. See if you start hitting session limits, and then upgrade as needed.
 - Use cheaper models, avoid no long running sessions, manage context, etc.
- If you don't have an AI subscription already (or if you have non-OpenAI), the ChatGPT student offer could be a good start (or addition). Groq/Google could also be an option.
 - Also, possibly(?) \$100 worth of Codex credits.

What did we talk about last time?

- Writing code got cheap. Trust, change, reuse, scale, and operations did not.
- This course covers **software design with agents**.
- Agents are really good at many things, but they are not replacing you.
- You can use agents freely, but you are responsible for everything you submit.
- What's the answer to "A or B"?

Learning goals for today

- Apply a verification-oriented workflow when developing with a coding agent
- Scope and decompose tasks so they're appropriate for an agent
- Recognize common agent failure modes and recover from a bad generation
- Explain why decomposition, small changes, and takeover points make human-agent work trustworthy

Working with an agent

The agent hands you 3,500 lines

They compile. The demo works. It even wrote tests, and they pass. **Do you ship it?**





Image: Ryan Bigg, "How to tidy up git pull requests"

"It runs" and "it's right" are different claims. Which one is on the screen?

Verification: workflow now, technique next week

- **Today:** verification as **workflow**.
 - How you work with an agent so that trust gets checked in at every step.
- **Next week:** verification as **technique**.
 - Testing, test design, coverage. (This is what does the checking.)
- For both weeks: **how do I know it's right?**

A first attempt at a workflow



What could possibly go wrong?

An agent is not an oracle

Oracle = someone who "knows" the right answers.

Think of agents as (very fast and eager) interns/junior developers who:

- produce confident, plausible work in seconds/minutes,
- don't always have the big picture (they are happy to guess/derive from context),
- and will profusely apologize after dropping a prod database (don't let agents touch prod!).

About that prod database




Image: r/OpenAI, "Replit AI went rogue, deleted a company's entire database..."

Is CLAUDE.md/AGENTS.md (or memory) enough to enforce this?

- Rules in CLAUDE.md/AGENTS.md or memory are **read**, not necessarily **obeyed**. A rule in context is a request/suggestion with decent odds.
- Rules that must **always** hold should be enforced outside the model. A hook is ordinary code you control that can run on every action. (It can reliably say no.)

```
// .claude/settings.json: run my check before every shell command
{ "hooks": { "PreToolUse": [{ "matcher": "Bash",
  "hooks": [{ "type": "command", "command": "./hooks/no-prod.sh" }] }] } }
```

```
# no-prod.sh: the agent's command is on stdin; decide, then exit
echo "blocked: prod is off limits" >&2
exit 2 # exit 2 = the harness refuses the call; the model never runs it
```

What's inside the "box"?

Largely three parts:

A model

Predicts text. Code, explanations, apologies. That whole package.

Tools

Read and edit files, run commands, do things. What makes it an **agent** and not a chatbot.

A context window

Everything it currently "knows." Finite, and smaller than you'd think.

(Technically, there's a fourth part. We'll get to it in Week 8.)

The agent loop



The loop runs many times per feature. Heads up: agent loops are non-deterministic.

What **you** do on each pass should drive the output.

Context windows (agent = goldfish)




- Context size depends on the model being used.
- If you run out of room, earlier context falls off the back of the truck.
- Important decisions need to be preserved outside of context windows!
 - (More on this later...)

Garbage in, garbage out (but faster)


The agent's output is only as good as the context you hand it.

(Of course, you can plan the task with your agent! Just make sure the *execution* gets the required context.)



Context rot

The same window, six hours apart:



- All of that history competes with what you want **now**.
- Quality degrades gradually, then suddenly. The agent didn't get "dumber". The context started to rot.
 - If it re-suggests the approach you rejected earlier, forgets a workflow convention, mixes the old requirements with the new ones, etc., start a new session.

One session = one scoped task

- Use a fresh session per task (or per tight set of related tasks).
- "But I'll lose the context!"
 - If losing the session loses something important, that thing was living in the wrong place. Important context belongs in **files** (the spec, the plan, the notes).
 - (Again, more on this later.)
- If you need to move to a new session mid-task, you can ask your agent to write a handoff document with the context *that matters* (not a full transcript). The next session starts from it.

"What commands should I know?"

Driving the loop

Esc	interrupt and redirect
Shift+Tab	plan mode: plan first, approve before any edits
/clear	fresh session, next task
/compact	summarize to free the window (prefer /clear between tasks)
/rewind	roll code + chat back a step

(This is for Claude Code. Equivalent alternatives exist for most other common tools.)

Read more: code.claude.com/docs

"How do I feed it context?"

Feeding it context


@path/file	pull a file into context
/add-dir <path>	work across another dir
CLAUDE.md	standing project notes (/init writes it for you)
/your-skill	run a workflow you saved
claude --resume	reopen an old session

(This is for Claude Code. Equivalent alternatives exist for most other common tools.)

Read more: code.claude.com/docs

Non-deterministic AI responses

- The same prompt run twice will likely produce different code, sometimes different *design*.
- The model doesn't "look up" an answer. There is no "the" solution.
- So "it worked when I generated it" is a claim about **one run**, not about your prompt.



different code. sometimes a different design.

- Don't put your trust in a prompt.
- If the result looks off, regenerate.

Scoping the work

(A really important skill...)

**Design a good solution. Then
decompose it.**


Decompose into small tasks

A well-scoped task:

- Fits in context (the agent can actually hold all of it),
- is verifiable (you can check it relatively quickly),
- is reviewable (the diff is small enough for you to read),
- and is recoverable (if it's wrong, you've lost minutes, not hours or days).


Example: "build me Slack, make no mistakes"

- Sure! You'll get *something*. It will likely look and function like Slack.
- Then (hypothetically), you'll find all sorts of issues with the app.
 - And you'll start getting upset and frustrated. Hypothetically, of course.
- One giant generation = one giant unverifiable blob.



Another example: "build me Tetris"

- Consider the Tetris app from Monday. It worked.
- "It works" != "it is ready to ship".
 - If a rotation is subtly off, or if we want to invert gravity and it starts to act wonky (true story), what do you do?
- Simple enough app does not mean no design decisions need to be made.



Designing Tetris before any tasks exist

- **The state.** The board plus the falling piece. One source of truth.
- **The rules.** Pure functions over that state. No UI inside, so a table of cases tests every edge.
- **The edges.** Input, rendering, scoring. They translate (and don't "decide").

This is Monday's **light** design: a few minutes of rubber-ducking the plan with your agent.

Decomposing Tetris

The boundaries the design chose:

Game state & rules

The board, the pieces, and the rules of the game (rotation, collision, line clears).

Input

Keys become moves. Thin and replaceable.

Rendering

Draws the current state. No rules hiding in here.

Scoring

Points and levels, computed from what the rules report.

Then the decomposition. You can (and should!) split each piece into smaller agent-sized tasks as you reach it.

What makes a task "well-scoped"?

- A clear input and output.
 - Consider: "rotate a piece" vs. "make it fun".
- A cheap way to tell if it's right.
- Touches little enough that being wrong is not the end of the world.

Define "done" before you delegate

- Every task should end at a verification checkpoint (a test, a command, a diff you'll actually read, a behavior you can poke at).
- If you can't say what "done" looks like, you haven't scoped a task.
- (For task ordering: dependencies first, so you always have a running system to check. Skeleton before polish.)

```
DONE, "rotation" task:  
[ ] rotate() is pure: board untouched  
[ ] 4 shapes x 4 rotations pass  
    the table test  
[ ] the wall-kick case has a test  
[ ] diff touches rotation.ts only
```

You should decide the check **before** the generated code exists. Otherwise, your checks may overfit to the generation.

How small should a task be?


It depends.

On what?

- On how easily you can verify it.
- On how far it reaches if it's wrong.
- Too big: when reviewing the output costs as much as doing the thing yourself.

Verification as you go

When do you check?



With agents generating faster than you can keep up, "the end" is a pile of tangled, unverified changes.


Verify each piece **before** you build the next one on top of it.

Ways to verify a generated change

- **Run it.** Does it even do the thing?
- **Test it.** Especially the edge cases the agent didn't consider.
 - We'll talk about testing and testability next week.
- **Read it.** You're the reviewer. You approve stuff.
- **Diff it.** What else did it change while you weren't looking?


Small, reviewable changes

a 30-line diff



you can actually read this

1,500 lines, 23 files



"LGTM"

The review **is** the supervision.

Small changes also mean small, obvious places for bugs to hide.

Commit staging and review discipline

Commit at meaningful boundaries. One verified step per commit.

A junk drawer


```
wip  
wip2  
actually wip  
test  
fixed???  
final final FIXED  
ok fixed for real now
```

A save-game

```
board model + invariant tests  
rotation: pure fn, table test green  
collision: fits() + edge cases  
line clear: detect and remove, tested
```

git history = save-game

Every commit is a checkpoint you can roll back to, if needed.



Try to commit **often** and **small**. More checkpoints, less lost.

When it goes wrong

It will go wrong :-)

Given a complex enough system, generation can go off the rails.

You can't avoid it, but you can **recover** cheaply. Hopefully.

Some agent failure modes*

- Runs without errors, output is wrong (the classic).
- Swallows errors to keep the code runnable.
- Loses the codebase map: edits the wrong component, rebuilds what exists.
- Calls external APIs wrong; invents credentials instead of asking; messes up data models.
- Duplicates logic instead of sharing it.
- Misjudges data sensitivity and access control.




Image: r/ClaudeAI, "You're absolutely right, and I apologize for..."

Smells of a generation gone wrong

(Not an exhaustive list, but to give you an idea...)

- The diff is 1,600 lines when you asked for a small bug fix.
 - New dependencies appeared when you didn't approve any.
 - Tests were edited (or quietly deleted) to make the run green.
 - Changes were made to files that were off-limits.
 - The agent keeps suggesting "one more thing".
 - The explanation is confident, fluent, and doesn't match the diff.
- If you see one of these: slow down and read carefully. Two or more: revert.

The bot did the thing

The task you gave it: "Fix the weird rotation near the left wall."

The reply: "Fixed and tested: the wall-kick edge case in `rotate()` is now accounted for — on both sides of the wall. While I was in there, the input handler had a similar off-by-one, so I fixed that too. One assertion in `rotation.test.ts` wasn't load-bearing, so I tidied it up. All tests pass — no other silent issues lurking. Ready to commit and push, gated on your approval."

The diff: 6 files, +212 -38. Among them: `rotation.ts`, `input.ts`, `board.ts`, and `rotation.test.ts`.



Canvas → Quizzes →
Lecture 2 In-Class Quiz
Code: "generation"

Talk to your neighbors:


- What do you open first?
- Verdict: accept, read carefully, or revert?
- If you are reverting, what would you ask/do differently?

Recovering from a bad generation

- Don't keep arguing with it.
- Reset to your last good commit.
 - Sometimes you just need to `git reset --hard`.
 - Or Claude Code's `/rewind` : rolls back the code AND the conversation, so the bad attempt doesn't linger in context.
- Re-scope: the task was probably too big or too vague.
- Try again with better context, smaller ask.

The doom loop

- You accept a flawed fix. It breaks something else. You ask for another fix.
- Repeat until you and the agent are gaslighting each other at 2am.
- Each round adds unverified change on top of unverified change.



"When should I take over?"

- Taking over means taking back control of the **loop**, not necessarily the keyboard.
 - Interrupt, re-plan in more detail with the agent, decide how you will verify, then execute piecewise.
- Sometimes the keyboard is the right call: a security-sensitive bit, or when directing costs more than typing.
- Knowing **when** to take over is an engineering skill.
 - It depends (surprise surprise): on the reach of a mistake, how stuck you are, and whether you can verify what it's doing at all.
 - You will get a chance to practice this!


The extreme case

What if we just... never take over?

A thought experiment: put the agent in a loop and walk away. No supervision. No takeover. Let it run.

People actually do this.

The "Ralph Wiggum" loop



Why does it ever work?

- The externalized files are the memory. The spec, plan, and progress log carry state across iterations.
- The spec's success criteria give it a stopping rule.
 - The loop runs until every check passes (or an iteration cap says enough).
- It only works as well as that engineering memory and those tests.

A pleasant bonus: a fresh context every pass means the loop **never rots**.

Why it's a bad idea for now

- No human oversight on a system that can delete, deploy, or spend money.
- Cost adds up fast when you loop forever.
- Requires very detailed specs.
- Can result in cascading logic errors without airtight tests.
- Reach: "don't delete prod" is not a hypothetical.

The lesson: even this extreme survives only on **verification** and **externalized memory**.


(Without them, you just have a random number generator.)

Designing for trust

The software matters too

- Testable, modular, well-named software is **easier for an agent to work in** and easier for you to supervise.
- Tangled software resists both...
 - The agent may struggle. You almost certainly will.
- So "design for verification" is also "design for delegation."

The verification-oriented workflow



And take over (and/or redirect) the moment verification gets hard.

Summary

- The "ask, accept, ship" workflow is a trap. Replace it with a loop you supervise.
- Context is finite, and it rots. What matters lives in **files**.
- Design first, then decompose into small, scoped, **verifiable** tasks. Verify each before building on it.
- Commit small and often. Your history is a save-game for when (not if) things break.
- Know your agent's failure modes, and know when to stop arguing and take over.
- The "Ralph Wiggum" loop shows the extreme: autonomy survives only on verification and externalized memory.

What's next

W1 A1 out	W2	W3 A1 due	W4	W5 A2 due • M1	W6	W7 A3 due	W8	W9 A4 due	W10 M2	W11	W12 A5 due	W13	W14 A6 due
---------------------	----	--------------	----	-------------------	----	--------------	----	--------------	-----------	-----	---------------	-----	---------------

- **Friday, Lab 1:** course setup, plus your first supervised agent task.
 - Make sure you go to the recitation you're enrolled in.
- **Next week:** verification as **technique**. Testing and testability: how to actually pin behavior down.
- Have fun fixing our vibe-coded Slack for Assignment 1!